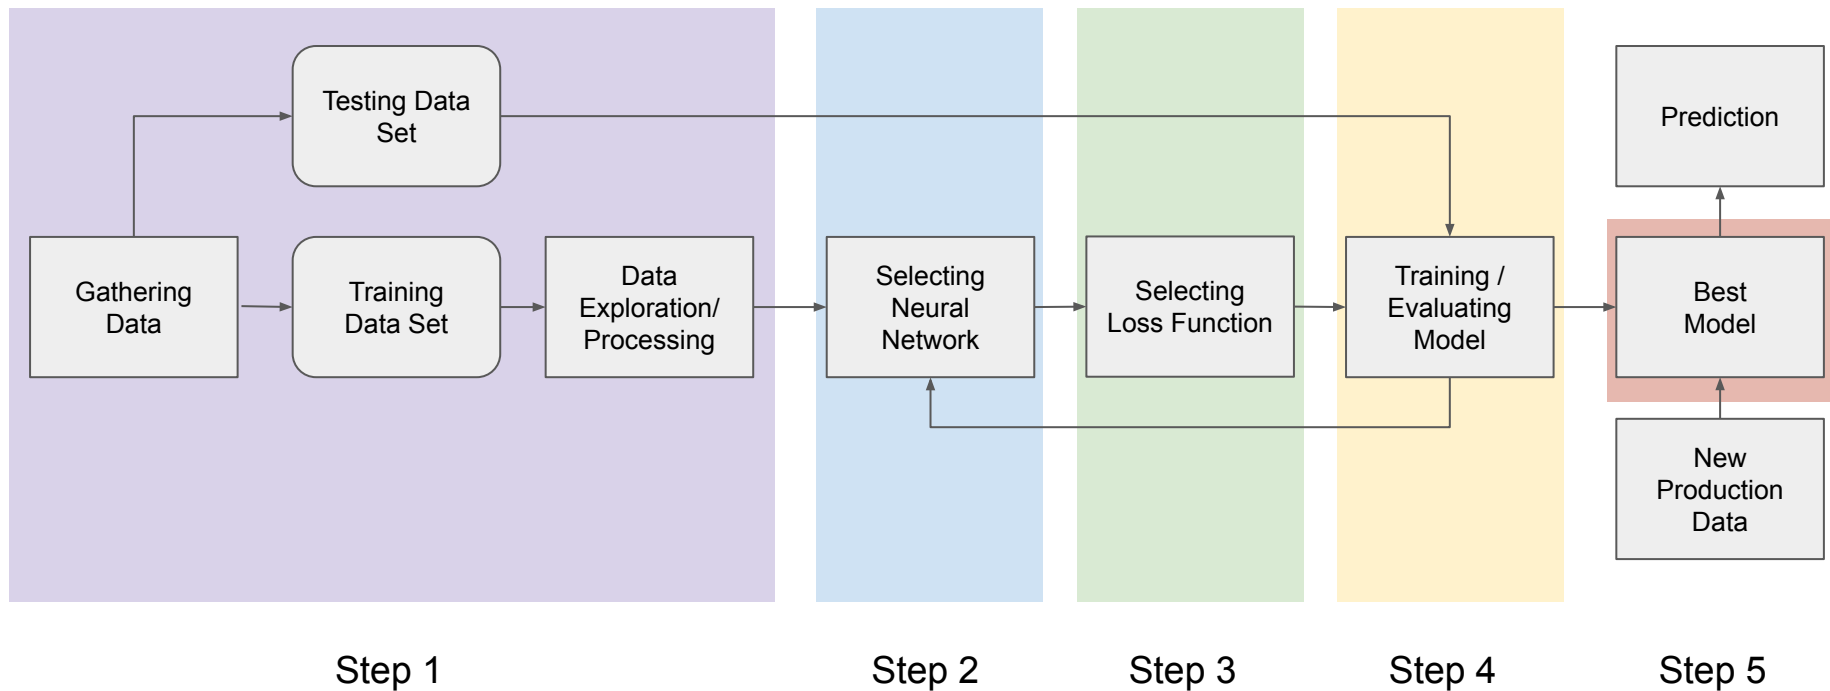


Learning Deep Learning with PyTorch

(2) Mechanics of Learning

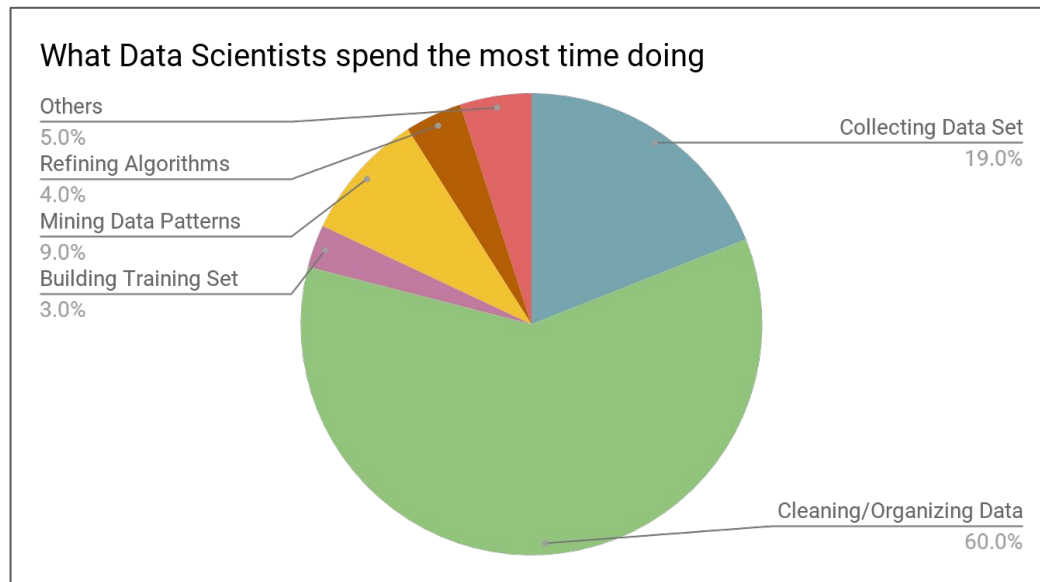
Qiyang Hu
UCLA IDRE
Oct 20, 2020

Workflow for a deep learning project



Step 1. Data Processing

- The most time-consuming but the most *creative* job
 - Take > 80% time
 - Require experience
 - May need domain expertise
- Determines the upper limit for the goodness of DL
 - Models/Algorithms: just approach the upper limit



Survey from Forbes in 2017 ([Data Source](#))

Typical tasks in the data processing step

- Data Preparation

- Gathering and cleaning
- Counting and statistics
- Annotation and ground truth labelling

- Data Tokenization

- Breaking the sequence data into units
- Mapping units to vectors
- Aligning & padding sequences

- Data Augmentation

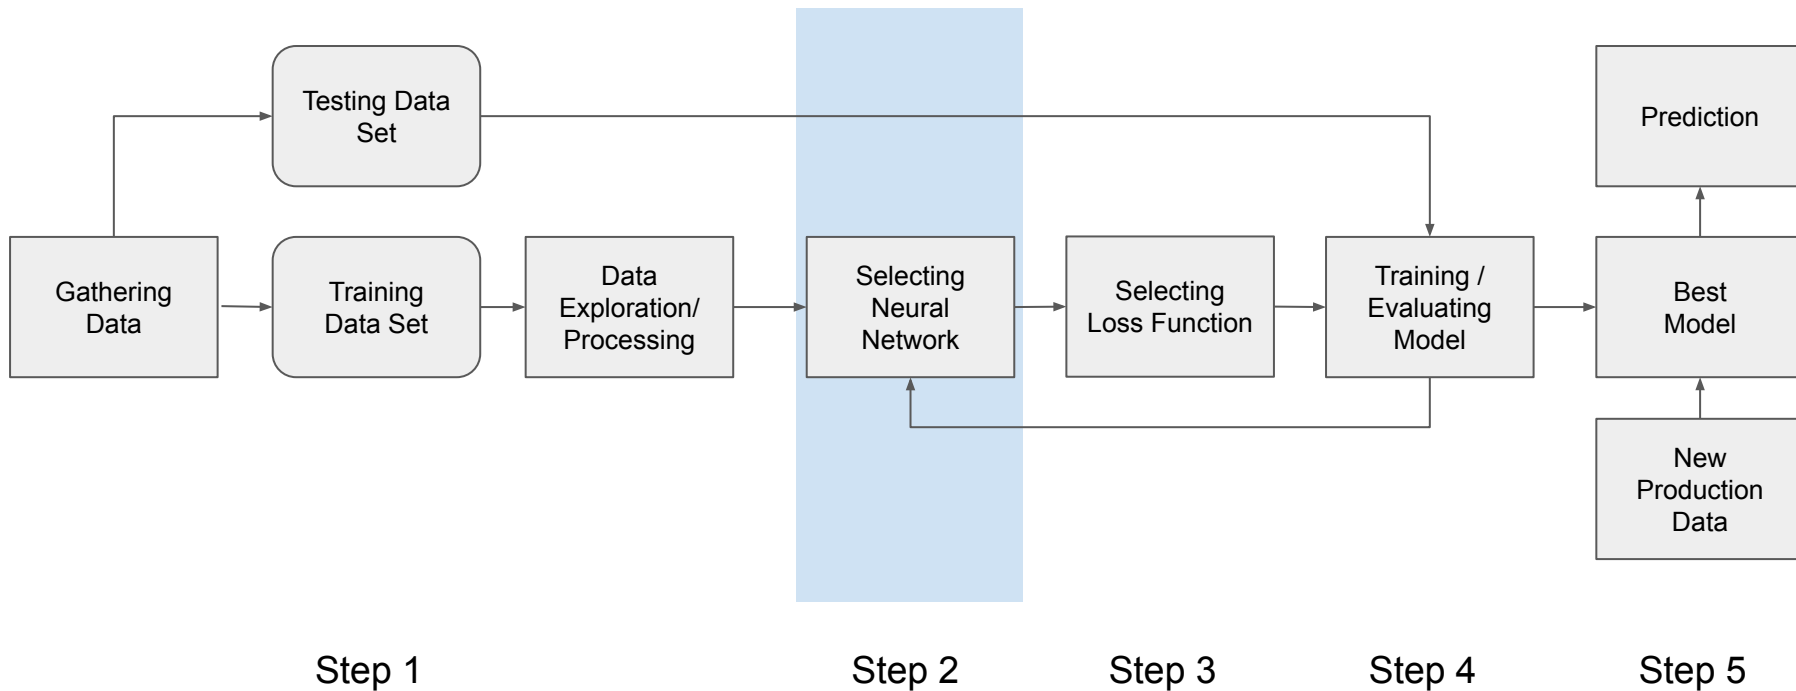
- Generate more training data

- Data Embedding

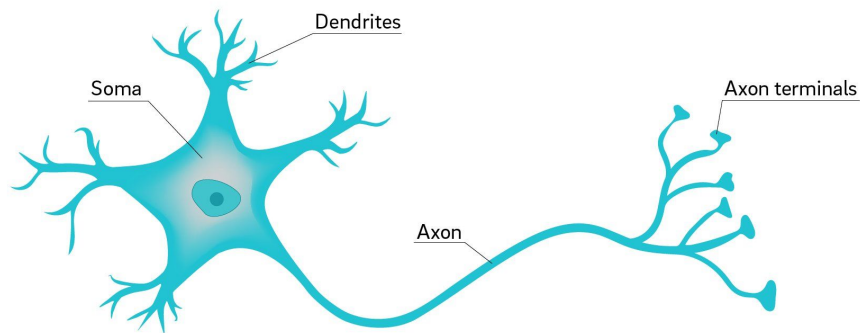

- Map data to lower-dim vectors
 - Sparse to dense
 - Merging diverse data
 - Preserve relationship
- Techniques
 - Std Dimensionality Reduction
 - Word2Vec
 - Be part of the model training
- *Representation Learning*

$$\text{Embedding Dims} \approx \sqrt[4]{\text{Possible Values}}$$


Workflow for a deep learning project



Recap: A linear classifier ~ one artificial neuron





(Deep) Neural Networks ~ piling/stacking logistic-regression classifiers




Why deep?

- Shallow network can fit any function
 - Has less number of hidden layers
 - Has to be really “fat”
- Deep network is more efficient.
 - It can extract/build better features
 - Exponentially fewer parameters ([2017](#))



Types of Deep Learning Architectures


Neural Networks



A higher-level classification of neural network types


- **Feed forward neural networks** (No cycle in node connections)

- Fully connected network
- Convolutional networks (CNNs)




- **Recurrent networks** (w/ directed cycle in node connections)

- Fully recurrent NN
- Recursive NN
- Long short-term memory (LSTM)
- Hopfield network (w/o hidden nodes)



- **Symmetric networks** (no directions in node connections)

- Boltzmann Machines
 - RBM, DBM



Activation Function

- Sigmoid function: $\sigma(z) = \frac{1}{1 + \exp(-z)}$
- tanh function: $\tanh(z) = \frac{\exp(z) - \exp(-z)}{\exp(z) + \exp(-z)}$

- Rectified linear unit (ReLU)

- Softplus
- Leaky ReLU
- Exponential LU (ELUs)
- GELU, etc.



$$f(x) = x^+ = \max(0, x)$$

- Softmax function:


$$y_i = \frac{e^{z^{(i)}}}{\sum_{j=0}^K e^{z^{(j)}}}$$

- Maxout Network:

- *Learnable* activation function



Workflow for a deep learning project



How to measure the performance of the model?

- General name: objective function
- Measure the misfit of the model as a function of parameters
 - Criterion is to *minimize* the error functions
 - Loss Function, Cost Function: a penalty on difference between predictions and labels
- Evaluate the probability of *generating* training set
 - Criterion is to *maximize* the distribution likelihood as a function of parameters
 - Maximum (log)-likelihood estimation: minimize the divergence of distributions
- Regression losses and classification losses

Loss functions

- Generative/Predictive:



- Regression Loss

- Mean Square Error / Quadratic Loss / L2 Loss:

$$L_{MSE} = \frac{1}{n} \sum_i^n (t_i - s_i)^2$$

- Mean Absolute Error / L1 Loss:

$$L_{MAE} = \frac{1}{n} \sum_i^n |t_i - s_i|$$

- Cross-Entropy Loss and variations

- Softmax Loss / Log Loss / Negative Log Likelihood

- Weighted CE / Balanced CE / Focal Loss

- Dice Loss / IOU Loss / Tversky Loss


$$L_{CE} = - \sum_i^C t_i \log(s_i)$$

- Contrastive:




- Ranking Loss/Margin Loss/Contrastive Loss/Triplet Loss

Workflow for a deep learning project





Training a DNN is an optimization problem



- We know how to compute $L(\theta)$, analytically or numerically.
- Start from an arbitrary initialization of θ_o , and get an initial $L_o(\theta)$
- Apply optimization algorithm to minimize $L(\theta)$

DL Optimization Algorithm

- Gradient Descent (a 1st-order approach) $\theta \leftarrow \theta - \eta \nabla L(\theta)$
 - Most popular algorithm
 - Pros: simple and fast
 - Cons: sometimes hard to tune



[Source Link](#)

Gradient-Descent Optimizers

- Stochastic GD / Mini-Batch GD
- Adding momentum:
 - Classical Momentum (CM)
 - Nesterov's Accelerated Gradient (NAG)
- Adaptive learning rate:
 - AdaGrad, AdaDelta, ...
 - RMSprop
- Combining the two
 - **ADAM** (as **default** in many libs)
- Beyond Adam:
 - Lookahead ([2019](#)), RAdam ([2019](#))
 - AdaBound/AmsBound ([ICLR 2019](#))
 - Range ([2019](#))
 - AdaBelief ([NeurIPS 2020 Spotlight](#))

Gradient descent vs Momentum vs
AdaGrad vs RMSProp vs Adam

[\(Source\)](#)

Higher Order Optimization Algorithm

- Newton-like methods (2nd-order methods)

$$\theta \leftarrow \theta - \frac{\ell'(\theta)}{\ell''(\theta)}$$

- Prons:
 - **Fewer** iterations (quadratic convergence)
 - Fewer hyperparameters
- Cons:
 - Much more **costly** in each iteration
 - Need more storing
- DFP/Broyden/BFGS/L-BFGS: a quasi-newton one
 - Good for low dimensional models
- Conjugate gradient (CG): between GD and Newton
 - moderately high dimensional models






Figure from [Wikipedia](https://en.wikipedia.org/wiki/Newton's_method)

Using Gradient Descent to train DNN



Backpropagation: a game of chain rule



$$y = \sigma_L \left(w_L \cdot \sigma_{L-1} \left(\cdots w_2 \cdot \sigma_1 \left(\underbrace{w_1 \cdot x + b_1}_{z_1} \right) + b_2 \right) + b_L \right)$$

$$\frac{\partial C(y(w) - \hat{y})}{\partial w} = \frac{\partial z}{\partial w} \frac{\partial C}{\partial z} = \frac{\partial z}{\partial w} \left[\frac{\partial a}{\partial z} \frac{\partial C}{\partial a} \right] = \frac{\partial z}{\partial w} \left[\sigma' \cdot \left(\underbrace{\frac{\partial z_{(+1)}}{\partial a} \frac{\partial C}{\partial z_{(+1)}}}_{a_1} \right) \right]$$

① Forward Pass


$$\frac{\partial z_1}{\partial w_1} = x_1 \longrightarrow \frac{\partial z_2}{\partial w_2} = a_1 \longrightarrow \cdots \longrightarrow \frac{\partial z_{L-1}}{\partial w_{L-1}} = a_{L-2} \longrightarrow \frac{\partial z_L}{\partial w_L} = a_{L-1}$$

② Backward Pass

$$\frac{\partial C}{\partial z_1} = \sigma'_1 \left[w_2 \frac{\partial C}{\partial z_2} \right] \longleftarrow \cdots \longleftarrow \frac{\partial C}{\partial z_{L-1}} = \sigma'_{L-1} \left[w_L \frac{\partial C}{\partial z_L} \right] \longleftarrow \frac{\partial C}{\partial z_L} = \sigma'_L \frac{\partial C}{\partial y} \longleftarrow \frac{\partial C}{\partial y}$$

Differentiability concerns on ReLU

- ReLU as one of the most popular activation functions: $f(x) = x^+ = \max(0, x)$
- ReLU is not differentiable at $x=0$
- Why we can use it in gradient based DNN training?
 - NN training *rarely* arrives at a local minimum of the cost function
 - Software implementations of NN training usually return one of the one-sided derivatives (*sub-gradient*)
- In practice we can safely disregard the non-differentiability of the hidden unit activation functions.



Workflow for a deep learning project

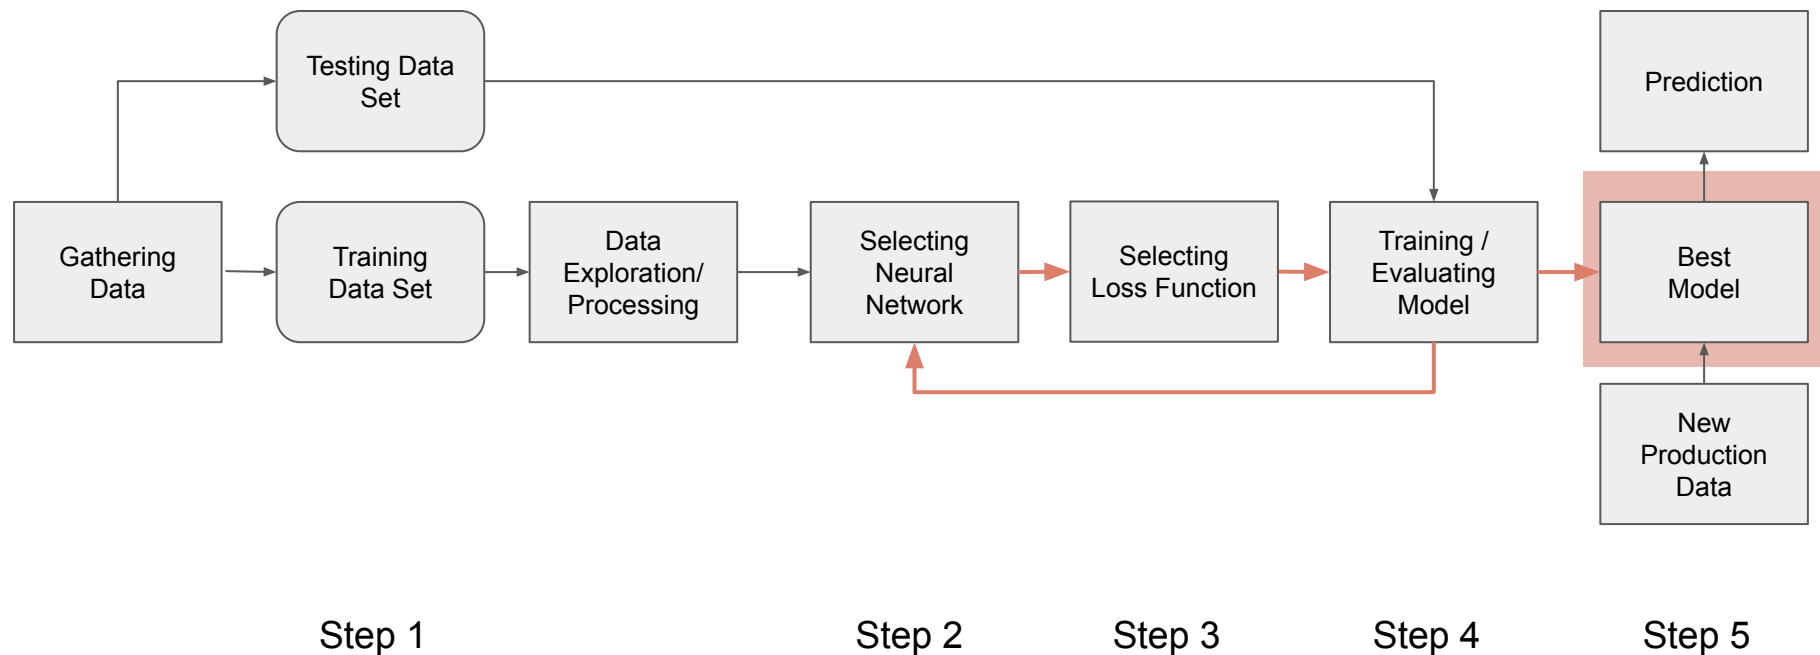




Figure [Source](#)



Underfitting and Overfitting

- Underfitting: model too simple:
 - Diagnose:
 - cannot even fit the training data
 - training error \sim testing error
 - Ignore the variance in training data
 - Higher prediction bias
- Overfitting: model too complex
 - Diagnose:
 - well-fit for training data
 - large error for testing data
 - Over-interpret training data
 - More deviation from new data



From Belkin's 2018 [paper](#)

How to prevent underfitting?


- Redesign the model
- Increase model's complexity
- Add more features as input
- Training longer
- More data will not help

How to prevent overfitting?

- Get more data
 - Collect more data
 - Data augmentation
- Reduce the model's complexity
- Regularization
 - Weight Regularization to make the model smoother (L1, L2, Elastic net)
$$\hat{L}(x, y) = L(x, y) + \lambda \sum_{i=1}^n \theta_i^2$$
 - Early stopping

Gradient vanishing/exploding in DL training

- Causes

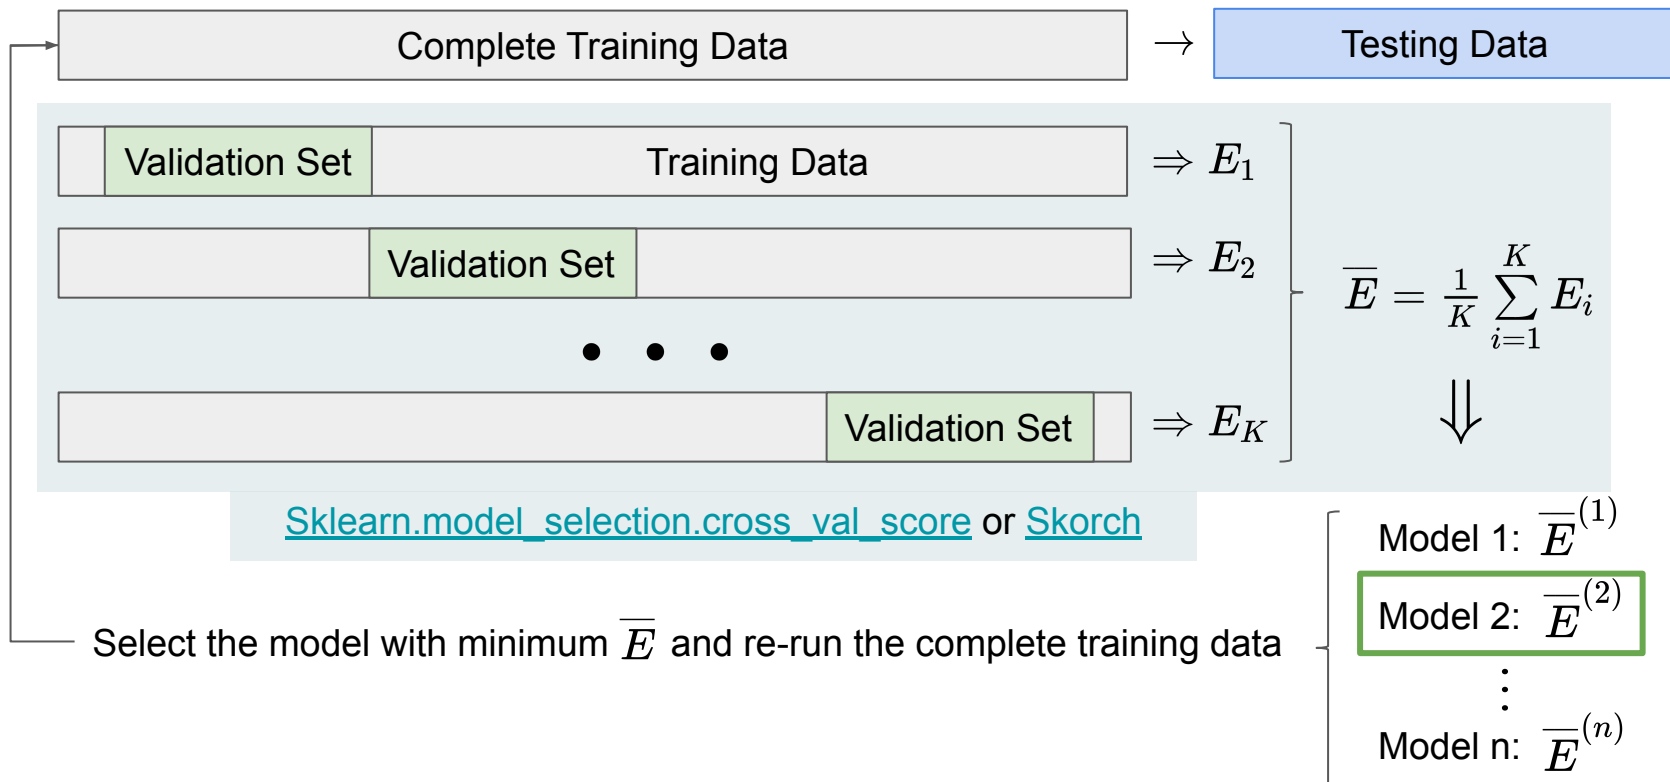


- Gradients in initial layers = Multiplication of Gradients at prior layers
- Small variation around 1 results in vanishing/exploding

- Techniques to resolve:

- General: adjusting learning rate, dropout, batch normalization, layer normalization
- For gradient exploding: gradient clipping, weight regularization
- For gradient vanishing: activation function, proper initialization parameters, LSTM, skip connections

Model Selection: K-fold Cross Validation



Errors/scores in practice



Error: E^{val} $<$ E^{Pub} $<$ E^{Pri}

Score: S^{val} $>$ S^{Pub} $>$ S^{Pri}

Don't forget to

- Github Repo:
 - <https://github.com/huqy/idre-learning-deep-learning-pytorch>
- Slack workspace:
 - bit.ly/Join-LDL
- Contact me
 - huqy@idre.ucla.edu
 - Direct message in Slack